

Data Wrangling with R's tidyverse

Wilfried Cools & Lara Stas

2023-12-13

square.research.vub.be

Compiled on R 4.3.1

! What-Why-Who

This site aims to introduce researchers to data manipulation in R with the `dplyr`, `tidyr`, and `stringr` packages of the **tidyverse** ecosystem.


Our target audience is primarily the research community at VUB / UZ Brussel, those who have some basic experience in R and want to know more.

We invite you to help improve this document by sending us feedback: wilfried.cools@vub.be

💡 Key Message

- Data manipulation can prepare data and/or its summarizing statistics
 - for modeling purposes
 - for visualisation purposes
- Data manipulation is inherent to data analysis, not just a precursor
 - no -fit's all data representation
 - * note: raw data should not be altered and kept safe
 - flexible use of data manipulation
 - * supports more informative and complete modeling
 - * elicits better visualisation of data and statistics
- Data manipulation is best done with coding
 - efficiently and correctly process data and statistics
 - maintain structure and transparency, to support reproducibility
- Data manipulation is easier and more intuitive when maintaining tidy data.

- tidy data: meaning appropriately mapped into structure
 - * each row an observation as research unit,
 - * each column a variable as property,
 - * each cell a particular value, linking row to column
 - * note: data can be split into multiple tables (relational data).
- aim for tidy data registration (avoid tedious manipulations)
- Workflow (Hadley Wickham):



R's tidyverse packages: dplyr and tidyr

- Current focus on **dplyr** and **tidyr** on manipulating and tidying data in the **tidyverse** eco-system (Hadley Wickham et al.)
- Data manipulation can be done in base R, or other packages
- **dplyr** and **tidyr**, the current defaults
 - inspired heavily on relational database logic
 - developed purposefully
 - * largely consistent
 - * well appreciated defaults
 - * easy and intuitive to build (if you get it)
 - * without losing much flexibility
- **dplyr** and **tidyr**, part of the **tidyverse** ecosystem includes:
 - **ggplot2** for visualizing data and statistics [check Visualization]
 - **stringr** for dealing with texts
 - **forcats** for dealing with factors
 - ...

Convenient cheat sheets at <https://rstudio.com/resources/cheatsheets/>.

Getting ahead of ourselves with dplyr

toy dataset

- The infamous `mtcars` data are used.
 - observe it's structure with `str()` and first 6 observations `head()` function.
 - note: available data with `data()`
- Have a tidyverse look at the data with `glimpse()`

```
glimpse(mtcars)
```

Rows: 32

Columns: 11

```
$ mpg <dbl> 21.0, 21.0, 22.8, 21.4, 18.7, 18.1, 14.3, 24.4, 22.8, 19.2, 17.8,~  
$ cyl <dbl> 6, 6, 4, 6, 8, 6, 8, 4, 4, 6, 6, 8, 8, 8, 8, 8, 8, 4, 4, 4, 4, 8,~  
$ disp <dbl> 160.0, 160.0, 108.0, 258.0, 360.0, 225.0, 360.0, 146.7, 140.8, 16~  
$ hp <dbl> 110, 110, 93, 110, 175, 105, 245, 62, 95, 123, 123, 180, 180, 180~  
$ drat <dbl> 3.90, 3.90, 3.85, 3.08, 3.15, 2.76, 3.21, 3.69, 3.92, 3.92, 3.92,~  
$ wt <dbl> 2.620, 2.875, 2.320, 3.215, 3.440, 3.460, 3.570, 3.190, 3.150, 3.~  
$ qsec <dbl> 16.46, 17.02, 18.61, 19.44, 17.02, 20.22, 15.84, 20.00, 22.90, 18~  
$ vs <dbl> 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0,~  
$ am <dbl> 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0,~  
$ gear <dbl> 4, 4, 4, 3, 3, 3, 3, 4, 4, 4, 4, 3, 3, 3, 3, 3, 3, 4, 4, 4, 3, 3,~  
$ carb <dbl> 4, 4, 1, 1, 2, 1, 4, 2, 2, 4, 4, 3, 3, 3, 4, 4, 4, 1, 2, 1, 1, 2,~
```

- Have a tidyverse look at the data with `slice_head()`

```
mtcars %>% slice_head(n=6)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

exemplary data manipulation

- Get the minimum value of the ratio of `mpg` over `hp` for each combination of `am` and `cyl`.
 - take the `mtcars` data,
 - select variables `mpg`, `cyl`, `hp`, `am`, and rename `hp` to `hpow`,
 - subset rows where `hpow` bigger than 3.5,
 - create new variable `mpgr` as the ratio `mpg` on `hpow`,
 - summarize `mpgr` as the minimum for every combination of `cyl` and `am`,
 - and reshape the result into a table with one row per `cyl`-value (4,6,8) and a column for each `am` value (0,1),
 - with column variable names renamed to `am0` and `am1`.

```
mtcars %>%  
  select(mpg, cyl, hpow=hp, am) %>%  
  filter(hpow > 3.5) %>%  
  mutate(mpgr = mpg/hpow) %>%  
  group_by(cyl, am) %>%  
  summarize(min=min(mpgr)) %>%  
  pivot_wider(names_from=am,  
              values_from=min) %>%  
  select(cyl, am0=`0`, am1=`1`)
```

cyl	am0	am1
4	0.2216495	0.1963303
6	0.1560976	0.1125714
8	0.1068571	NA

dplyr package, functions to manipulate data

- `dplyr` reflects the `apply` function in base R
 - `d` is for data frames
- Focus on manipulating data frames (tibbles):
 - subsetting, altering, summarizing, ordering, combining, reshaping
- The main -verbs- (see example above)

- `filter()` : conditional selection of cases
- `select()` : conditional selection of variables, allows reordering and renaming
- `mutate()` : creation of new variables based on existing variables
- `summarise()` : reduce sets of values to single values
- The verb to structure data (see example above)
 - `group_by()` : internal grouping, undo with `ungroup()`
 - works preceding main verbs
- The verbs to enhance control on scope (advanced)
 - `across()` : new way of scoping (instead of `*_it`, `*_at`, `*_all`)
 - * works for selection in `mutate()` and `summarize()`
- Additional `dplyr` verbs:
 - `arrange` : ordering of cases
 - `sample_n` and `sample_frac` : random sampling
 - `slice`, `transmute`, `rename`, `relocate`, ...
- Verbs to extend data
 - `bind_rows` and `bind_cols` : append data of same structure
 - `left_join`, `right_join`, `inner_join`, `full_join`, `semi_join` and `anti_join` : join data using indicator variable(s)
- Note: only the core of `dplyr` is discussed, much more is possible

group_by()

- Grouping prepares data for group specific operations

intro

- Get a glimpse of the data as before,
 - number of rows and columns
 - * in tidy data: observations and variables
 - number of groups, and grouping variables
 - * 4 groups: 2 am x 2 vs
 - * Note: grouping structure part of `glimpse`-output
- The width is set for presentation purposes

```
tst <- mtcars %>% group_by(am,vs)
glimpse(tst,width=40)
```

```
Rows: 32
Columns: 11
Groups: am, vs [4]
$ mpg <dbl> 21.0, 21.0, 22.8, 21.4, 1~
$ cyl <dbl> 6, 6, 4, 6, 8, 6, 8, 4, 4~
$ disp <dbl> 160.0, 160.0, 108.0, 258.~
$ hp <dbl> 110, 110, 93, 110, 175, 1~
$ drat <dbl> 3.90, 3.90, 3.85, 3.08, 3~
$ wt <dbl> 2.620, 2.875, 2.320, 3.21~
$ qsec <dbl> 16.46, 17.02, 18.61, 19.4~
$ vs <dbl> 0, 0, 1, 1, 0, 1, 0, 1, 1~
$ am <dbl> 1, 1, 1, 0, 0, 0, 0, 0, 0~
$ gear <dbl> 4, 4, 4, 3, 3, 3, 3, 4, 4~
$ carb <dbl> 4, 4, 1, 1, 2, 1, 4, 2, 2~
```

- Actions on grouped data are grouped too,
 - e.g., a frequency table, count the number of observations (`count()`)
 - grouped data result in grouped counts

```
tst %>% count()
```

am	vs	n
0	0	12
0	1	7
1	0	6
1	1	7

- Remove grouping with `ungroup()`
 - this is good practice to avoid unwanted effects !

```
tst <- tst %>% ungroup( )
tst %>% count()
```

n

- Alternatively, overwrite the initial grouping
 - the last grouping is used by default
 - additional arguments, for example `.add` and `.drop`, can change that
 - * a first groups by `vs`
 - * a second groups by `am` and `vs`

```
mtcars %>% group_by(am) %>% group_by(vs)
mtcars %>% group_by(am) %>% group_by(vs, .add=TRUE)
```

- Transformed variables can also be used for grouping
 - e.g., cutting the `mpg` in 3 groups with `cut()` then use `count()`
Notice the intervals that are created.

```
tst <- mtcars %>% group_by(mpg3 = cut(mpg, 3))
tst %>% count()
```

mpg3	n
<10.4,18.2]	14
<18.2,26.1]	13
<26.1,33.9]	5

exercises

- Embedded within the next sections

filter()

- Filtering returns rows using matching conditions

intro

- Get a subset of rows that includes only those rows with `mpg` above 30

```
mtcars %>% filter(mpg > 30)
```

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2

- Include more than just one condition,
 - take only rows with `mpg` above 20 AND `qsec` below or equal to 18
 - note: consecutive filtering achieves the same.
 - * & for **and**
 - * for **or**
 - * ! for **not**

```
mtcars %>% filter(mpg > 20 & qsec <= 18)
```

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2

- More complex conditions can be specified
 - take rows with `mpg` above 30 OR `qsec` below 20 AND `am` equal to 0
 - all the rules of logic apply, parentheses included

```
mtcars %>% filter(mpg > 30 | (qsec > 20 & am==0))
```

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2

33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2

- Grouping also works here
 - get all `distinct` values of `cyl` per level of `gear`
 - note: this selects the first unique rows

```
mtcars %>% group_by(gear) %>% distinct(cyl)
```

gear	cyl
4	6
4	4
3	6
3	8
3	4
5	4
5	8
5	6

exercises

- The `starwars` dataset is already part of tidyverse, so you should have it available !

```
data(starwars)
```

- Have a glimpse at the data, what data types are included ?

```
glimpse(starwars)
```

```
Rows: 87
Columns: 14
$ name      <chr> "Luke Skywalker", "C-3PO", "R2-D2", "Darth Vader", "Leia Or~
$ height    <int> 172, 167, 96, 202, 150, 178, 165, 97, 183, 182, 188, 180, 2~
$ mass      <dbl> 77.0, 75.0, 32.0, 136.0, 49.0, 120.0, 75.0, 32.0, 84.0, 77.~
$ hair_color <chr> "blond", NA, NA, "none", "brown", "brown, grey", "brown", N~
$ skin_color <chr> "fair", "gold", "white, blue", "white", "light", "light", "~
$ eye_color  <chr> "blue", "yellow", "red", "yellow", "brown", "blue", "blue",~
$ birth_year <dbl> 19.0, 112.0, 33.0, 41.9, 19.0, 52.0, 47.0, NA, 24.0, 57.0, ~
```

```

$ sex      <chr> "male", "none", "none", "male", "female", "male", "female", ~
$ gender   <chr> "masculine", "masculine", "masculine", "masculine", "femini~
$ homeworld <chr> "Tatooine", "Tatooine", "Naboo", "Tatooine", "Alderaan", "T~
$ species  <chr> "Human", "Droid", "Droid", "Human", "Human", "Human", "Huma~
$ films    <list> <"The Empire Strikes Back", "Revenge of the Sith", "Return~
$ vehicles <list> <"Snowspeeder", "Imperial Speeder Bike">, <>, <>, <>, "Imp~
$ starships <list> <"X-wing", "Imperial shuttle">, <>, <>, "TIE Advanced x1",~

```

- Do note, different data types are included in the tibble (data frame)
 - chr for characters, int for integers, dbl for doubles, we miss the lgl for a boolean
 - notice that even a vector of type list can be included.
- Filter the rows to subset the data and retain only characters with light skin and brown eye color

```
starwars %>% filter(skin_color == "light", eye_color == "brown")
```

```

# A tibble: 7 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Leia Org~    150    49 brown      light      brown          19 fema~ femin~
2 Biggs Da~    183    84 black      light      brown          24 male  mascu~
3 Cordé       157    NA brown      light      brown          NA fema~ femin~
4 Dormé       165    NA brown      light      brown          NA fema~ femin~
5 Raymus A~    188    79 brown      light      brown          NA male  mascu~
6 Poe Dame~    NA     NA brown      light      brown          NA male  mascu~
7 Padmé Am~    165    45 brown      light      brown          46 fema~ femin~
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

- Arrange the data according the character's height, largest on top ! (google it!!)

```
starwars %>% arrange(desc(height))
```

```

# A tibble: 87 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Yarael ~    264    NA none      white      yellow          NA male  mascu~
2 Tarfful    234   136 brown      brown      blue          NA male  mascu~
3 Lama Su    229    88 none      grey      black          NA male  mascu~
4 Chewbac~    228   112 brown      unknown    blue          200 male  mascu~

```

```

5 Roos Ta~      224      82 none      grey      orange      NA      male      mascu~
6 Grievous      216     159 none      brown, wh~ green, y~    NA      male      mascu~
7 Taun We       213      NA none      grey      black      NA      fema~    femin~
8 Rugor N~      206      NA none      green      orange      NA      male      mascu~
9 Tion Me~      206      80 none      grey      black      NA      male      mascu~
10 Darth V~     202     136 none      white      yellow      41.9    male      mascu~
# i 77 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

- Who is smallest (comes on top after arranging) ?

```
starwars %>% arrange(height)
```

```

# A tibble: 87 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Yoda        66    17 white      green      brown        896 male  mascu~
2 Ratts T~     79    15 none      grey, blue unknown      NA male  mascu~
3 Wicket ~     88    20 brown      brown      brown         8 male  mascu~
4 Dud Bolt     94    45 none      blue, grey yellow      NA male  mascu~
5 R2-D2        96    32 <NA>      white, bl~ red        33 none  mascu~
6 R4-P17        96    NA none      silver, r~ red, blue   NA none  femin~
7 R5-D4        97    32 <NA>      white, red red        NA none  mascu~
8 Sebulba     112    40 none      grey, red  orange      NA male  mascu~
9 Gasgano     122    NA none      white, bl~ black      NA male  mascu~
10 Watto      137    NA black      blue, grey yellow      NA male  mascu~
# i 77 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

- Slice the data and keep only the 5th to 10th observation ! (?slice)

```
starwars %>% slice(5:10)
```

```

# A tibble: 6 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Leia Org~     150    49 brown      light      brown        19 fema~  femin~
2 Owen Lars     178   120 brown, gr~ light      blue         52 male  mascu~

```

```

3 Beru Whi~    165    75 brown    light    blue              47 fema~ femin~
4 R5-D4        97     32 <NA>    white, red red        NA none  mascu~
5 Biggs Da~    183    84 black    light    brown            24 male  mascu~
6 Obi-Wan ~    182    77 auburn, w~ fair    blue-gray        57 male  mascu~
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

- Slice the first 2 observations for each gender (group your data) !
 - what other functions are discussed at ?`slice_head` ?

```
starwars %>% group_by(gender) %>% slice_head(n=2)
```

```

# A tibble: 6 x 14
# Groups:   gender [3]
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Leia Org~   150    49 brown    light    brown          19 fema~ femin~
2 Beru Whi~   165    75 brown    light    blue           47 fema~ femin~
3 Luke Sky~   172    77 blond    fair     blue           19 male  mascu~
4 C-3PO      167    75 <NA>    gold     yellow        112 none  mascu~
5 Ric Olié   183    NA brown    fair     blue           NA <NA> <NA>
6 Quarsh P~  183    NA black    dark     brown          62 <NA> <NA>
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

- Use `slice_sample()` to randomly select 5 observations !

```
starwars %>% slice_sample(n = 5)
```

```

# A tibble: 5 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Owen Lars   178   120 brown, gr~ light    blue          52 male  mascu~
2 Mace Win~   188    84 none     dark     brown          72 male  mascu~
3 Ki-Adi-M~   198    82 white    pale     yellow         92 male  mascu~
4 Beru Whi~   165    75 brown    light    blue           47 fema~ femin~
5 Bail Pre~   191    NA black    tan      brown          67 male  mascu~
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

- Use `slice_max()` to select 3 observations with highest values on `height` !

```
starwars %>% slice_max(height, n = 3)
```

```
# A tibble: 3 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Yarael P~    264    NA none      white      yellow      NA male masculi~
2 Tarfful      234   136 brown     brown      blue        NA male masculi~
3 Lama Su      229    88 none      grey       black       NA male masculi~
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>
```

- Get the top 3 (highest mass) for each species !
 - ignore characters with missing data for `mass`
 - note, -not missing- are those who are not ! missing `is.na()`

```
starwars %>% group_by(species) %>% filter(!is.na(mass)) %>% slice_max(mass, n = 3)
```

```
# A tibble: 40 x 14
# Groups:   species [32]
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Ratts T~     79    15 none      grey, blue unknown      NA male masculi~
2 Dexter ~    198   102 none      brown      yellow      NA male masculi~
3 Ki-Adi-~    198    82 white     pale       yellow      92 male masculi~
4 Zam Wes~    168    55 blonde   fair, gre~ yellow      NA fema~ femin~
5 IG-88       200   140 none      metal      red        15 none masculi~
6 C-3P0       167    75 <NA>      gold       yellow     112 none masculi~
7 R2-D2        96    32 <NA>      white, bl~ red        33 none masculi~
8 R5-D4        97    32 <NA>      white, red red        NA none masculi~
9 Sebulba     112    40 none      grey, red  orange     NA male masculi~
10 Wicket ~     88    20 brown     brown      brown       8 male masculi~
# i 30 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>
```

`select()`

- Extract columns (variables) by name (or position), rename and/or reorder them

intro

- Select the variable `mpg`
 - notice that even with one column, the result remains a dataframe (not a vector), this is tidyverse policy !
- An operation on a data with a certain type should result in data of the same type.
 - if you take one column from a matrix you have a one column matrix, not a vector.
 - if you take one column from a data frame, again, you end up with a one-column data frame, not a vector.

```
mtcars %>% select(mpg)
```

mpg
21.0
21.0
22.8
21.4
18.7
18.1

- To retrieve a vector with `dplyr` use `pull()`
- Specific operations allow for changing the data types

```
mtcars %>% pull(mpg)
```

```
[1] 21.0 21.0 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 17.8 16.4 17.3 15.2 10.4  
[16] 10.4 14.7 32.4 30.4 33.9 21.5 15.5 15.2 13.3 19.2 27.3 26.0 30.4 15.8 19.7  
[31] 15.0 21.4
```

- Extract columns `qsec` and `mpg` (top 6 observations)
 - note: more than one column can be considered jointly, their order is specified as such

```
mtcars %>% select(qsec,mpg)
```

qsec	mpg
16.46	21.0
17.02	21.0
18.61	22.8
19.44	21.4
17.02	18.7
20.22	18.1

- Extract the third and first column (top 6)
 - note: columns can be extracted by their position

```
mtcars %>% select(3,1)
```

disp	mpg
160	21.0
160	21.0
108	22.8
258	21.4
360	18.7
225	18.1

- Remove columns at third to sixth position (top 6)
 - note: to remove, use a negation, but it is either keep or remove not both

```
mtcars %>% select(-c(3:6))
```

mpg	cyl	qsec	vs	am	gear	carb
21.0	6	16.46	0	1	4	4
21.0	6	17.02	0	1	4	4
22.8	4	18.61	1	1	4	1
21.4	6	19.44	1	0	3	1
18.7	8	17.02	0	0	3	2
18.1	6	20.22	1	0	3	1

- **helper functions** can facilitate selections
- Use partial string matching with `contains()`

- extract columns with names that include the string `ar` (show 6)

```
mtcars %>% select(contains('ar'))
```

gear	carb
4	4
4	4
4	1
3	1
3	2
3	1

- Use regular expressions with `matches()`
 - extract columns with names that include the string `ar` but with at least one element before and after it

```
mtcars %>% select(matches('.ar.'))
```

carb
4
4
1
1
2
1

- Variables can be renamed during selection
 - rename the `cyl` into `cyl468` to reflect its values
 - same for `vs` and `am`, and select it together with `mpg`

```
mtcars %>% select(mpg, cyl468=cyl, vs01=vs, am01=am)
```

mpg	cyl468	vs01	am01
21.0	6	0	1
21.0	6	0	1
22.8	4	1	1
21.4	6	1	0

18.7	8	0	0
18.1	6	1	0

- Rename the `cyl`, `vs` and `am` directly

```
mtcars %>% rename(cyl468=cyl,vs01=vs,am01=am)
```

mpg	cyl468	disp	hp	drat	wt	qsec	vs01	am01	gear	carb
21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

- Note that a `select()` will include the grouping variables by default
- Grouping variables are identified with `group_cols()`
- Create a grouping by `vs` and `am`, and extract only those columns

```
mtcars %>% group_by(vs,am) %>% select(group_cols( ))
```

0 - 0
0 - 1
1 - 0
1 - 1

exercises

- The `starwars` dataset is probably still loaded into your workspace !

```
data(starwars)
```

- Select the columns `hair`, `skin` and `eye color` !

```
starwars %>% select(hair_color, skin_color, eye_color)
```

```
# A tibble: 87 x 3
  hair_color skin_color eye_color
  <chr>      <chr>      <chr>
1 blond     fair        blue
2 <NA>      gold        yellow
3 <NA>      white, blue red
4 none      white       yellow
5 brown     light       brown
6 brown, grey light       blue
7 brown     light       blue
8 <NA>      white, red  red
9 black     light       brown
10 auburn, white fair       blue-gray
# i 77 more rows
```

- Use the : operator for consecutive columns hair and eye color !

```
starwars %>% select(hair_color:eye_color)
```

```
# A tibble: 87 x 3
  hair_color skin_color eye_color
  <chr>      <chr>      <chr>
1 blond     fair        blue
2 <NA>      gold        yellow
3 <NA>      white, blue red
4 none      white       yellow
5 brown     light       brown
6 brown, grey light       blue
7 brown     light       blue
8 <NA>      white, red  red
9 black     light       brown
10 auburn, white fair       blue-gray
# i 77 more rows
```

- Remove these columns instead of selecting them !

```
starwars %>% select(-(hair_color:eye_color))
```

```
# A tibble: 87 x 11
  name      height  mass birth_year sex  gender homeworld species films vehicles
  <chr>      <int> <dbl>      <dbl> <chr> <chr>  <chr>    <chr>  <lis> <list>
1 Luke S~    172    77        19  male mascul~ Tatooine Human  <chr> <chr>
2 C-3P0      167    75       112  none mascul~ Tatooine Droid  <chr> <chr>
3 R2-D2       96    32        33  none mascul~ Naboo   Droid  <chr> <chr>
4 Darth ~    202   136       41.9 male mascul~ Tatooine Human  <chr> <chr>
5 Leia O~    150    49        19  fema~ femin~ Alderaan Human  <chr> <chr>
6 Owen L~    178   120        52  male mascul~ Tatooine Human  <chr> <chr>
7 Beru W~    165    75        47  fema~ femin~ Tatooine Human  <chr> <chr>
8 R5-D4       97    32        NA  none mascul~ Tatooine Droid  <chr> <chr>
9 Biggs ~    183    84        24  male mascul~ Tatooine Human  <chr> <chr>
10 Obi-Wa~    182    77        57  male mascul~ Stewjon Human  <chr> <chr>
# i 77 more rows
# i 1 more variable: starships <list>
```

- Select all columns with a name ending with `color` (check help files on helper functions, use `?language`) !

```
starwars %>% select(ends_with("color"))
```

```
# A tibble: 87 x 3
  hair_color  skin_color eye_color
  <chr>      <chr>      <chr>
1 blond      fair        blue
2 <NA>       gold        yellow
3 <NA>       white, blue red
4 none       white       yellow
5 brown      light       brown
6 brown, grey light       blue
7 brown      light       blue
8 <NA>       white, red  red
9 black      light       brown
10 auburn, white fair       blue-gray
# i 77 more rows
```

- Use `select` to rename `homeworld` to `home_world` !

```
starwars %>% select(home_world = homeworld)
```

```
# A tibble: 87 x 1
  home_world
  <chr>
1 Tatooine
2 Tatooine
3 Naboo
4 Tatooine
5 Alderaan
6 Tatooine
7 Tatooine
8 Tatooine
9 Tatooine
10 Stewjon
# i 77 more rows
```

- Do the same with the `rename()` function !

```
starwars %>% rename(home_world = homeworld)
```

```
# A tibble: 87 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>    <int> <dbl> <chr>    <chr>    <chr>    <dbl> <chr> <chr>
1 Luke Sk~    172    77 blond    fair      blue      19    male masculin
2 C-3PO     167    75 <NA>     gold      yellow    112   none masculin
3 R2-D2      96    32 <NA>     white, bl~ red       33    none masculin
4 Darth V~   202   136 none     white     yellow    41.9  male masculin
5 Leia Or~   150    49 brown    light     brown     19    fema~ feminin
6 Owen La~   178   120 brown, gr~ light     blue     52    male masculin
7 Beru Wh~   165    75 brown    light     blue     47    fema~ feminin
8 R5-D4      97    32 <NA>     white, red red       NA    none masculin
9 Biggs D~   183    84 black    light     brown     24    male masculin
10 Obi-Wan~   182    77 auburn, w~ fair      blue-gray  57    male masculin
# i 77 more rows
# i 5 more variables: home_world <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>
```

- Select only the numeric variables, use `where()` and `is.numeric()` !
 - Maybe check the `?language` again

```
starwars %>% select(where(is.numeric))
```

```
# A tibble: 87 x 3
  height mass birth_year
  <int> <dbl>    <dbl>
1    172    77        19
2    167    75       112
3     96    32        33
4    202   136       41.9
5    150    49        19
6    178   120        52
7    165    75        47
8     97    32        NA
9    183    84        24
10   182    77        57
# i 77 more rows
```

- Select only those variables with names `height`, `mass` and/or `size`, with `any_of()` !

```
starwars %>% select(any_of(c('height','mass','size')))
```

```
# A tibble: 87 x 2
  height mass
  <int> <dbl>
1    172    77
2    167    75
3     96    32
4    202   136
5    150    49
6    178   120
7    165    75
8     97    32
9    183    84
10   182    77
# i 77 more rows
```

mutate()

- Create new variables based on existing ones

intro

- A new variable (column) `mpg2` can be created by `mpg` value squared

```
mtcars %>% mutate(mpg2=mpg^2)
```

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	mpg2
21.0	6	160	110	3.90	2.620	16.46	0	1	4	4	441.00
21.0	6	160	110	3.90	2.875	17.02	0	1	4	4	441.00
22.8	4	108	93	3.85	2.320	18.61	1	1	4	1	519.84
21.4	6	258	110	3.08	3.215	19.44	1	0	3	1	457.96
18.7	8	360	175	3.15	3.440	17.02	0	0	3	2	349.69
18.1	6	225	105	2.76	3.460	20.22	1	0	3	1	327.61

- The original value can also be overwritten
 - e.g., the `mpg` can be assigned the values of `mpg` squared

```
mtcars %>% mutate(mpg=mpg^2)
```

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
441.00	6	160	110	3.90	2.620	16.46	0	1	4	4
441.00	6	160	110	3.90	2.875	17.02	0	1	4	4
519.84	4	108	93	3.85	2.320	18.61	1	1	4	1
457.96	6	258	110	3.08	3.215	19.44	1	0	3	1
349.69	8	360	175	3.15	3.440	17.02	0	0	3	2
327.61	6	225	105	2.76	3.460	20.22	1	0	3	1

- Based on multiple variables, e.g., `NEWVAR` can represent the `mpg` value multiplied by the `vs` value
- The convenient `everything` function is a short-cut to every column not explicitly mentioned

```
mtcars %>% mutate(NEWVAR=mpg*vs) %>% select(NEWVAR,everything())
```

NEWVAR	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0.0	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
0.0	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
22.8	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1

21.4	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
0.0	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
18.1	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

- A new variable can be created based on a newly created variable as well
 - e.g., NEWVAR is the mpg value multiplied by the vs value and this new variable is divided by the disp variable

```
mtcars %>% mutate(NEWVAR=mpg*vs,NEWVAR2=NEWVAR/disp) %>% select(NEWVAR,NEWVAR2,everything())
```

NEWVAR	NEWVAR2	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0.0	0.00000000	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
0.0	0.00000000	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
22.8	0.21111111	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
21.4	0.08294574	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
0.0	0.00000000	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
18.1	0.08044444	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

- **window functions** facilitate the automation of mutations (google for dplyr window functions).
 - e.g., add a column with the cumulative sum of mpg using `cumsum()`

```
mtcars %>% mutate(NEWVAR=cumsum(mpg))
```

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	NEWVAR
21.0	6	160	110	3.90	2.620	16.46	0	1	4	4	21.0
21.0	6	160	110	3.90	2.875	17.02	0	1	4	4	42.0
22.8	4	108	93	3.85	2.320	18.61	1	1	4	1	64.8
21.4	6	258	110	3.08	3.215	19.44	1	0	3	1	86.2
18.7	8	360	175	3.15	3.440	17.02	0	0	3	2	104.9
18.1	6	225	105	2.76	3.460	20.22	1	0	3	1	123.0

- Add a column with indicator whether the mpg is between 20 and 22

```
mtcars %>% mutate(NEWVAR=between(mpg,20,22)) %>% select(NEWVAR,everything())
```

NEWVAR	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
TRUE	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
TRUE	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
FALSE	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
TRUE	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
FALSE	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
FALSE	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

- Add a row number dependent on the rank of `mpg` values, with the `rownumber()` function. When arranged by `mpg` this is more clear.

```
mtcars %>% mutate(id=row_number(mpg)) %>% select(id,everything())
mtcars %>% mutate(id=row_number(mpg)) %>% arrange(mpg) %>% select(id,everything())
```

id	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	10.4	8	472	205	2.93	5.250	17.98	0	0	3	4
2	10.4	8	460	215	3.00	5.424	17.82	0	0	3	4
3	13.3	8	350	245	3.73	3.840	15.41	0	0	3	4
4	14.3	8	360	245	3.21	3.570	15.84	0	0	3	4
5	14.7	8	440	230	3.23	5.345	17.42	0	0	3	4
6	15.0	8	301	335	3.54	3.570	14.60	0	1	5	8

- Grouping variables can group the operations
- To create a ranking within groups, `vs` and `am`, `row_number()` can be used again
 - Notice, for each combination of `vs` and `am`, there will be a 1 (first), 2 (second)... for `id` (not all shown, only 6 per combination are shown).

```
mtcars %>% group_by(vs,am) %>% mutate(id=row_number(mpg))
```

mpg	cyl	disp	hp	drat	wt	qsec	gear	carb	id
0 - 0									
18.7	8	360.0	175	3.15	3.440	17.02	3	2	11
14.3	8	360.0	245	3.21	3.570	15.84	3	4	4
16.4	8	275.8	180	3.07	4.070	17.40	3	3	9
17.3	8	275.8	180	3.07	3.730	17.60	3	3	10
15.2	8	275.8	180	3.07	3.780	18.00	3	3	6

10.4	8	472.0	205	2.93	5.250	17.98	3	4	1
0 - 1									
21.0	6	160.0	110	3.90	2.620	16.46	4	4	4
21.0	6	160.0	110	3.90	2.875	17.02	4	4	5
26.0	4	120.3	91	4.43	2.140	16.70	5	2	6
15.8	8	351.0	264	4.22	3.170	14.50	5	4	2
19.7	6	145.0	175	3.62	2.770	15.50	5	6	3
15.0	8	301.0	335	3.54	3.570	14.60	5	8	1
1 - 0									
21.4	6	258.0	110	3.08	3.215	19.44	3	1	4
18.1	6	225.0	105	2.76	3.460	20.22	3	1	2
24.4	4	146.7	62	3.69	3.190	20.00	4	2	7
22.8	4	140.8	95	3.92	3.150	22.90	4	2	6
19.2	6	167.6	123	3.92	3.440	18.30	4	4	3
17.8	6	167.6	123	3.92	3.440	18.90	4	4	1
1 - 1									
22.8	4	108.0	93	3.85	2.320	18.61	4	1	2
32.4	4	78.7	66	4.08	2.200	19.47	4	1	6
30.4	4	75.7	52	4.93	1.615	18.52	4	2	4
33.9	4	71.1	65	4.22	1.835	19.90	4	1	7
27.3	4	79.0	66	4.08	1.935	18.90	4	1	3
30.4	4	95.1	113	3.77	1.513	16.90	5	2	5

exercises

- For the starwars data, create a new variable `height_m` with `height` divided by 100 !

```
starwars %>% mutate(height_m = height / 100) %>% select(height_m, height, everything( ))
```

```
# A tibble: 87 x 15
```

	height_m	height	name	mass	hair_color	skin_color	eye_color	birth_year	sex
	<dbl>	<int>	<chr>	<dbl>	<chr>	<chr>	<chr>	<dbl>	<chr>
1	1.72	172	Luke ~	77	blond	fair	blue	19	male
2	1.67	167	C-3P0	75	<NA>	gold	yellow	112	none
3	0.96	96	R2-D2	32	<NA>	white, bl~	red	33	none
4	2.02	202	Darth~	136	none	white	yellow	41.9	male
5	1.5	150	Leia ~	49	brown	light	brown	19	fema~
6	1.78	178	Owen ~	120	brown, gr~	light	blue	52	male

```

7      1.65      165 Beru ~      75 brown      light      blue      47 fema~
8      0.97       97 R5-D4      32 <NA>      white, red red      NA none
9      1.83      183 Biggs~     84 black      light      brown      24 male
10     1.82      182 Obi-W~     77 auburn, w~ fair      blue-gray    57 male
# i 77 more rows
# i 6 more variables: gender <chr>, homeworld <chr>, species <chr>,
#   films <list>, vehicles <list>, starships <list>

```

- Note, the `select` function also defines the order, again the `everything` function avoids explicitly naming all variables
- Create the same new variable, but also define BMI as `mass / height_m` to the power 2 !

```
starwars %>% mutate(height_m = height / 100, BMI = mass / (height_m^2)) %>% select(BMI, ev
```

```

# A tibble: 87 x 16
      BMI name      height  mass hair_color skin_color eye_color birth_year sex
  <dbl> <chr>    <int> <dbl> <chr>    <chr>    <chr>    <dbl> <chr>
1  26.0 Luke Sky~    172    77 blond     fair     blue      19 male
2  26.9 C-3PO      167    75 <NA>     gold     yellow    112 none
3  34.7 R2-D2       96    32 <NA>     white, bl~ red      33 none
4  33.3 Darth Va~   202   136 none     white     yellow    41.9 male
5  21.8 Leia Org~   150    49 brown     light     brown     19 fema~
6  37.9 Owen Lars  178   120 brown, gr~ light     blue     52 male
7  27.5 Beru Whi~   165    75 brown     light     blue     47 fema~
8  34.0 R5-D4       97    32 <NA>     white, red red     NA none
9  25.1 Biggs Da~   183    84 black     light     brown     24 male
10 23.2 Obi-Wan ~   182    77 auburn, w~ fair     blue-gray  57 male
# i 77 more rows
# i 7 more variables: gender <chr>, homeworld <chr>, species <chr>,
#   films <list>, vehicles <list>, starships <list>, height_m <dbl>

```

- Use `transmute` to repeat the above mutation but keep only `height_m` and `BMI` !

```
starwars %>% transmute(height_m = height / 100, BMI = mass / (height_m^2))
```

```

# A tibble: 87 x 2
  height_m BMI
  <dbl> <dbl>
1     1.72 26.0
2     1.67 26.9

```

```

3      0.96  34.7
4      2.02  33.3
5      1.5   21.8
6      1.78  37.9
7      1.65  27.5
8      0.97  34.0
9      1.83  25.1
10     1.82  23.2
# i 77 more rows

```

- Create a new variable with the z-score of height ($zscore = (value - mean) / sd$) !

```
starwars %>% mutate(zscore=(height-mean(height,na.rm=T))/sd(height,na.rm=T)) %>% select(zs
```

```
# A tibble: 87 x 15
```

	zscore	name	height	mass	hair_color	skin_color	eye_color	birth_year	sex
	<dbl>	<chr>	<int>	<dbl>	<chr>	<chr>	<chr>	<dbl>	<chr>
1	-0.0678	Luke S~	172	77	blond	fair	blue	19	male
2	-0.212	C-3PO	167	75	<NA>	gold	yellow	112	none
3	-2.25	R2-D2	96	32	<NA>	white, bl~	red	33	none
4	0.795	Darth ~	202	136	none	white	yellow	41.9	male
5	-0.701	Leia O~	150	49	brown	light	brown	19	fema~
6	0.105	Owen L~	178	120	brown, gr~	light	blue	52	male
7	-0.269	Beru W~	165	75	brown	light	blue	47	fema~
8	-2.22	R5-D4	97	32	<NA>	white, red	red	NA	none
9	0.249	Biggs ~	183	84	black	light	brown	24	male
10	0.220	Obi-Wa~	182	77	auburn, w~	fair	blue-gray	57	male

```

# i 77 more rows
# i 6 more variables: gender <chr>, homeworld <chr>, species <chr>,
#   films <list>, vehicles <list>, starships <list>

```

- Now create that z-score per species !

```
starwars %>% group_by(species) %>% mutate(zscore=(height-mean(height,na.rm=T))/sd(height,n
```

```
# A tibble: 87 x 15
```

```
# Groups:   species [38]
```

	zscore	species	height	name	mass	hair_color	skin_color	eye_color	birth_year
	<dbl>	<chr>	<int>	<chr>	<dbl>	<chr>	<chr>	<chr>	<dbl>
1	-0.371	Human	172	Luke ~	77	blond	fair	blue	19

```

2  0.728 Droid      167 C-3PO      75 <NA>      gold      yellow      112
3 -0.716 Droid      96 R2-D2      32 <NA>      white, bl~ red      33
4  2.02  Human     202 Darth~   136 none     white      yellow     41.9
5 -2.13  Human     150 Leia ~   49 brown    light      brown      19
6  0.108 Human     178 Owen ~   120 brown, gr~ light      blue      52
7 -0.929 Human     165 Beru ~   75 brown    light      blue      47
8 -0.696 Droid      97 R5-D4      32 <NA>      white, red red      NA
9  0.507 Human     183 Biggs~   84 black     light      brown      24
10 0.427 Human     182 Obi-W~   77 auburn, w~ fair      blue-gray   57
# i 77 more rows
# i 6 more variables: sex <chr>, gender <chr>, homeworld <chr>, films <list>,
#   vehicles <list>, starships <list>

```

- Create a gender indicator that replaces the male and female labels with m and f (use `recode()`)!

```
starwars %>% mutate(new_value=recode(sex,'male'='m','female'='f')) %>% select(new_value, s
```

```

# A tibble: 87 x 15
  new_value sex   name height mass hair_color skin_color eye_color birth_year
  <chr>      <chr> <chr>  <int> <dbl> <chr>      <chr>      <chr>      <dbl>
1 m         male   Luke~   172    77 blond     fair       blue       19
2 none      none   C-3PO   167    75 <NA>      gold       yellow     112
3 none      none   R2-D2   96     32 <NA>      white, bl~ red       33
4 m         male   Dart~   202   136 none     white      yellow     41.9
5 f         fema~   Leia~   150    49 brown    light      brown      19
6 m         male   Owen~   178   120 brown, gr~ light      blue       52
7 f         fema~   Beru~   165    75 brown    light      blue       47
8 none      none   R5-D4   97     32 <NA>      white, red red       NA
9 m         male   Bigg~   183    84 black     light      brown      24
10 m        male   Obi-~   182    77 auburn, w~ fair      blue-gray   57
# i 77 more rows
# i 6 more variables: gender <chr>, homeworld <chr>, species <chr>,
#   films <list>, vehicles <list>, starships <list>

```

- Create a gender indicator that, when sex is 'none' uses the species values and otherwise keeps the sex specification (use `ifelse()`)!

```
starwars %>% mutate(new_value=ifelse(sex=='none',species,sex)) %>% select(new_value, speci
```

```
# A tibble: 87 x 15
  new_value species sex    name    height  mass hair_color skin_color eye_color
  <chr>      <chr>  <chr>  <chr>    <int> <dbl> <chr>      <chr>    <chr>
1 male      Human   male   Luke S~   172    77 blond      fair      blue
2 Droid     Droid   none   C-3P0    167    75 <NA>      gold      yellow
3 Droid     Droid   none   R2-D2     96    32 <NA>      white, bl~ red
4 male      Human   male   Darth ~   202   136 none       white     yellow
5 female    Human   female Leia O~   150    49 brown      light     brown
6 male      Human   male   Owen L~   178   120 brown, gr~ light     blue
7 female    Human   female Beru W~   165    75 brown      light     blue
8 Droid     Droid   none   R5-D4     97    32 <NA>      white, red red
9 male      Human   male   Biggs ~   183    84 black      light     brown
10 male     Human   male   Obi-Wa~   182    77 auburn, w~ fair      blue-gray
# i 77 more rows
# i 6 more variables: birth_year <dbl>, gender <chr>, homeworld <chr>,
#   films <list>, vehicles <list>, starships <list>
```

summarize()

- Reduce sets of values into their summaries, based on grouped data.

intro

- A new variable (column) is created based on an existing one by summarizing, condensing the data
 - e.g., the mean of all `mpg` values can be obtained

```
mtcars %>% summarize(myAverage=mean(mpg))
```

myAverage
20.09062

- Multiple summaries can be obtained jointly, the mean and standard deviation of all `mpg` values can be obtained
 - you could do that for multiple variables, and also include `disp`

```
mtcars %>% summarize(myAvMpg=mean(mpg),mySdMpg=sd(mpg),myAvDisp=mean(dis),mySdDisp=sd(d
```

myAvMpg	mySdMpg	myAvDisp	mySdDisp
20.09062	6.026948	230.7219	123.9387

- Grouping variables are very natural to use with `summarize()`
- The mean of all `mpg` values can be obtained for each level of `vs`

```
mtcars %>% group_by(vs) %>% summarize(myAverage=mean(mpg))
```

vs	myAverage
0	16.61667
1	24.55714

- The mean and standard deviation can be obtained for multiple variables too,
 - e.g., an average and standard deviation of `mpg` and `disp` for each group

```
mtcars %>% group_by(vs,am) %>% summarize(myAvMpg=mean(mpg),mySdMpg=sd(mpg),myAvDisp=mean(d
```

am	myAvMpg	mySdMpg	myAvDisp	mySdDisp
0				
0	15.05000	2.774396	357.6167	71.82349
1	19.75000	4.008865	206.2167	95.23362
1				
0	20.74286	2.471071	175.1143	49.13072
1	28.37143	4.757701	89.8000	18.80213

- The total number of observations within a group
 - e.g., `vs`, can be obtained with `n()`, or using the special verb `count()`

```
mtcars %>% group_by(vs) %>% count( )
mtcars %>% group_by(vs) %>% summarize(mycount=n( ))
```

vs	mycount
0	18
1	14

- Making use of **summary functions**, summarizing can be more automated
 - e.g., the number of distinct values in a vector for each combination of **vs** and **am** can be obtained with `n_distinct()`, and the third number of each group with `nth()`

```
mtcars %>% group_by(vs,am) %>% summarize(nrDist=n_distinct(mpg),`3th`=nth(mpg,3))
```

am	nrDist	3th
0		
0	10	16.4
1	5	26.0
1		
0	7	24.4
1	6	30.4

exercises

- Summarize the **height** into the average height (some missing values need to be dealt with, check `?mean`) !

```
starwars %>% summarise(height = mean(height, na.rm = TRUE))
```

```
# A tibble: 1 x 1
  height
  <dbl>
1  174.
```

- Repeat the above, but include the average **mass** per **species** and **sex** !

```
starwars %>% group_by(species,sex) %>% summarize(height = mean(height, na.rm = TRUE), mass
```

```
# A tibble: 41 x 4
# Groups:   species [38]
  species    sex  height  mass
  <chr>     <chr>   <dbl> <dbl>
1 Aleena   male      79    15
2 Besalisk male     198   102
3 Cerean   male     198    82
4 Chagrian male     196   NaN
5 Clawdite female   168    55
6 Droid    none     131.  69.8
7 Dug      male     112    40
8 Ewok     male      88    20
9 Geonosian male     183    80
10 Gungan  male     209.   74
# i 31 more rows
```

across()

- Scoping a verb, the `across()` function allows for `summarize()` or `mutate()` operations on a set of variables

intro

- Select variables by either explicitly naming them or by extraction using dedicated functions
- Turn to the `mtcars` data again, take a `glimpse` to remind yourself what the data looks like
- Select `mpg`, `cyl`, `am` and `vs`
- Turn both `am` and `vs` into a factor before calling the structure with `glimpse()`

```
mtcars %>% select(mpg,cyl,am,vs) %>% mutate(across(c('am','vs'),factor)) %>% glimpse( )
```

Rows: 32

Columns: 4

```
$ mpg <dbl> 21.0, 21.0, 22.8, 21.4, 18.7, 18.1, 14.3, 24.4, 22.8, 19.2, 17.8, ~
$ cyl <dbl> 6, 6, 4, 6, 8, 6, 8, 4, 4, 6, 6, 8, 8, 8, 8, 8, 8, 4, 4, 4, 4, 8, ~
```



```
$ am <fct> 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, ~
$ vs <fct> 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, ~
```

- A factor is made from all variables in between `cyl` and `vs` with a `:` operator

```
mtcars %>% select(mpg,cyl,am,vs) %>% mutate(across(cyl:vs,factor)) %>% glimpse( )
```

Rows: 32

Columns: 4

```
$ mpg <dbl> 21.0, 21.0, 22.8, 21.4, 18.7, 18.1, 14.3, 24.4, 22.8, 19.2, 17.8, ~
$ cyl <fct> 6, 6, 4, 6, 8, 6, 8, 4, 4, 6, 6, 8, 8, 8, 8, 8, 8, 4, 4, 4, 4, 8, ~
$ am <fct> 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, ~
$ vs <fct> 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, ~
```

- A factor is made from all variables that contain the letter combination `ar`

```
mtcars %>% select(mpg,cyl,gear,carb) %>% mutate(across(contains("ar"),factor)) %>% glimpse
```

Rows: 32

Columns: 4

```
$ mpg <dbl> 21.0, 21.0, 22.8, 21.4, 18.7, 18.1, 14.3, 24.4, 22.8, 19.2, 17.8, ~
$ cyl <dbl> 6, 6, 4, 6, 8, 6, 8, 4, 4, 6, 6, 8, 8, 8, 8, 8, 8, 4, 4, 4, 4, 8, ~
$ gear <fct> 4, 4, 4, 3, 3, 3, 3, 4, 4, 4, 4, 3, 3, 3, 3, 3, 3, 4, 4, 4, 3, 3, ~
$ carb <fct> 4, 4, 1, 1, 2, 1, 4, 2, 2, 4, 4, 3, 3, 3, 4, 4, 4, 1, 2, 1, 1, 2, ~
```

- The `across()` function allows for applying a list of functions
 - e.g., for the first and third variable, a function is applied to obtain a median, a mean and an sd
 - use is made of `tidyverse` short-cuts `~` to indicate a function is used and `.x` that works as a container for the variables used in that function.

```
descr <- list(
  md = ~median(.x, na.rm = TRUE),
  av = ~mean(.x, na.rm = TRUE),
  sd = ~sd(.x, na.rm = TRUE)
)
mtcars %>% mutate(across(c(1,3), descr))
```

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	mpg_md	mpg_av	mpg_sd	disp_
21.0	6	160	110	3.90	2.620	16.46	0	1	4	4	19.2	20.09062	6.026948	1
21.0	6	160	110	3.90	2.875	17.02	0	1	4	4	19.2	20.09062	6.026948	1
22.8	4	108	93	3.85	2.320	18.61	1	1	4	1	19.2	20.09062	6.026948	1
21.4	6	258	110	3.08	3.215	19.44	1	0	3	1	19.2	20.09062	6.026948	1
18.7	8	360	175	3.15	3.440	17.02	0	0	3	2	19.2	20.09062	6.026948	1
18.1	6	225	105	2.76	3.460	20.22	1	0	3	1	19.2	20.09062	6.026948	1

- Making use of **helper functions**
 - the same as for `select()`, selections can be more automated
- Helper functions include among others `all_of()`, `where()`, `matches()`, `starts_with()`
 - are possible to use within `mutate()` and `summarize()`

```
descr <- list(
  md = ~median(.x, na.rm = TRUE),
  av = ~mean(.x, na.rm = TRUE),
  sd = ~sd(.x, na.rm = TRUE)
)
mtcars %>% mutate(across(c(1,3), descr))
```

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	mpg_md	mpg_av	mpg_sd	disp_
21.0	6	160	110	3.90	2.620	16.46	0	1	4	4	19.2	20.09062	6.026948	1
21.0	6	160	110	3.90	2.875	17.02	0	1	4	4	19.2	20.09062	6.026948	1
22.8	4	108	93	3.85	2.320	18.61	1	1	4	1	19.2	20.09062	6.026948	1
21.4	6	258	110	3.08	3.215	19.44	1	0	3	1	19.2	20.09062	6.026948	1
18.7	8	360	175	3.15	3.440	17.02	0	0	3	2	19.2	20.09062	6.026948	1
18.1	6	225	105	2.76	3.460	20.22	1	0	3	1	19.2	20.09062	6.026948	1

exercises

- For starwars data, request the minimum and maximum values of the numeric variables (use `where()`).
 - Note that you need to deal with missing values

```

min_max <- list(
  min = ~min(.x, na.rm = TRUE),
  max = ~max(.x, na.rm = TRUE)
)
starwars %>% summarise(across(where(is.numeric), min_max))

# A tibble: 1 x 6
  height_min height_max mass_min mass_max birth_year_min birth_year_max
    <int>      <int>    <dbl>    <dbl>         <dbl>         <dbl>
1       66      264      15     1358             8          896

```

join()

- Datafiles can be combined using common variables that serve as key (cfr. relational databases).

intro

- Methods differ primarily in how they deal with mismatches in key variable values
- Assume a cylinder specific datafile, `mtcyl`, with a 2 cylinder but no 8 cylinder unlike the `mtcars` (4,6,8)

```

mtcyl <- tribble(
  ~cyl,~type,
  2,'small',
  4,'medium',
  6,'large'
)

```

- Combine the `mtcars` and `mtcyl` but ignore the irrelevant `cyl` equal to 2 (not part of `mtcars`), with a `left_join()`
 - Notice that `cyl` equal to 8 turns out missing, because it is not specified in the -right- datafile (`mtcyl`)

```

mtcars %>% left_join(mtcyl) %>% select(cyl,type,everything())

```

Joining with 'by = join_by(cyl)'

cyl	type	mpg	disp	hp	drat	wt	qsec	vs	am	gear	carb
6	large	21.0	160	110	3.90	2.620	16.46	0	1	4	4
6	large	21.0	160	110	3.90	2.875	17.02	0	1	4	4
4	medium	22.8	108	93	3.85	2.320	18.61	1	1	4	1
6	large	21.4	258	110	3.08	3.215	19.44	1	0	3	1
8	NA	18.7	360	175	3.15	3.440	17.02	0	0	3	2
6	large	18.1	225	105	2.76	3.460	20.22	1	0	3	1

- Combine the `mtcars` and `mtcyl` but ignore the `cyl` equal to 8 because it lacks information on `type`, with a `right_join()`
 - Notice that `cyl` equal to 2 is included, but turns out missing for most variables because it is not specified in the -left- datafile

```
mtcars %>% right_join(mtcyl) %>% arrange(cyl) %>% select(cyl,type,everything())
```

Joining with `'by = join_by(cyl)'`

cyl	type	mpg	disp	hp	drat	wt	qsec	vs	am	gear	carb
2	small	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
4	medium	22.8	108.0	93	3.85	2.320	18.61	1	1	4	1
4	medium	24.4	146.7	62	3.69	3.190	20.00	1	0	4	2
4	medium	22.8	140.8	95	3.92	3.150	22.90	1	0	4	2
4	medium	32.4	78.7	66	4.08	2.200	19.47	1	1	4	1
4	medium	30.4	75.7	52	4.93	1.615	18.52	1	1	4	2

- Combine the `mtcars` and `mtcyl` for only those observations with the linking variable `cyl` in both files, with a `inner_join()`
 - Notice no missing values, but some data is not included

```
mtcars %>% inner_join(mtcyl) %>% arrange(cyl) %>% select(cyl,type,everything())
```

Joining with `'by = join_by(cyl)'`

cyl	type	mpg	disp	hp	drat	wt	qsec	vs	am	gear	carb
4	medium	22.8	108.0	93	3.85	2.320	18.61	1	1	4	1

4	medium	24.4	146.7	62	3.69	3.190	20.00	1	0	4	2
4	medium	22.8	140.8	95	3.92	3.150	22.90	1	0	4	2
4	medium	32.4	78.7	66	4.08	2.200	19.47	1	1	4	1
4	medium	30.4	75.7	52	4.93	1.615	18.52	1	1	4	2
4	medium	33.9	71.1	65	4.22	1.835	19.90	1	1	4	1

- Combine the `mtcars` and `mtcyl` keeping all available information, with a `full_join()` showing selected rows 1 to 3, 5, 7 and 33

```
mtcars %>% full_join(mtcyl) %>% slice(c(1:3,5,7,33))
```

Joining with `'by = join_by(cyl)'`

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	type
21.0	6	160	110	3.90	2.620	16.46	0	1	4	4	large
21.0	6	160	110	3.90	2.875	17.02	0	1	4	4	large
22.8	4	108	93	3.85	2.320	18.61	1	1	4	1	medium
18.7	8	360	175	3.15	3.440	17.02	0	0	3	2	NA
14.3	8	360	245	3.21	3.570	15.84	0	0	3	4	NA
NA	2	NA	NA	NA	NA	NA	NA	NA	NA	NA	small

- Other types of join exist, like `semi_join()`, `nest_join()`, `anti_join()`, which are described in the help files.

exercises

- Two mini tibbles `band_members` and `band_instruments` are probably loaded into your workspace automatically as part of the tidyverse !

```
band_members
```

```
# A tibble: 3 x 2
  name band
  <chr> <chr>
1 Mick  Stones
2 John  Beatles
3 Paul  Beatles
```

```
band_instruments
```

```
# A tibble: 3 x 2
  name plays
  <chr> <chr>
1 John guitar
2 Paul bass
3 Keith guitar
```

- Combine the two, left/right/inner/full !

```
band_members %>% inner_join(band_instruments)
```

Joining with 'by = join_by(name)'

```
# A tibble: 2 x 3
  name band plays
  <chr> <chr> <chr>
1 John Beatles guitar
2 Paul Beatles bass
```

```
band_members %>% left_join(band_instruments)
```

Joining with 'by = join_by(name)'

```
# A tibble: 3 x 3
  name band plays
  <chr> <chr> <chr>
1 Mick Stones <NA>
2 John Beatles guitar
3 Paul Beatles bass
```

```
band_members %>% right_join(band_instruments)
```

Joining with 'by = join_by(name)'

```
# A tibble: 3 x 3
  name band plays
  <chr> <chr> <chr>
1 John Beatles guitar
2 Paul Beatles bass
3 Keith <NA> guitar
```

```
band_members %>% full_join(band_instruments)
```

Joining with 'by = join_by(name)'

```
# A tibble: 4 x 3
  name band plays
  <chr> <chr> <chr>
1 Mick Stones <NA>
2 John Beatles guitar
3 Paul Beatles bass
4 Keith <NA> guitar
```

- Try out the same with `semi_join()` and `anti_join()` and interpret what happens !

```
band_members %>% semi_join(band_instruments)
```

Joining with 'by = join_by(name)'

```
# A tibble: 2 x 2
  name band
  <chr> <chr>
1 John Beatles
2 Paul Beatles
```

```
band_members %>% anti_join(band_instruments)
```

Joining with 'by = join_by(name)'

```
# A tibble: 1 x 2
  name band
  <chr> <chr>
1 Mick Stones
```

dplyr exercises, catching up

- Compare the structure of the `mtcars` data with a glimpse at that data.

```
glimpse(mtcars)
```

```
Rows: 32
Columns: 11
$ mpg <dbl> 21.0, 21.0, 22.8, 21.4, 18.7, 18.1, 14.3, 24.4, 22.8, 19.2, 17.8,~
$ cyl <dbl> 6, 6, 4, 6, 8, 6, 8, 4, 4, 6, 6, 8, 8, 8, 8, 8, 4, 4, 4, 4, 8,~
$ disp <dbl> 160.0, 160.0, 108.0, 258.0, 360.0, 225.0, 360.0, 146.7, 140.8, 16~
$ hp <dbl> 110, 110, 93, 110, 175, 105, 245, 62, 95, 123, 123, 180, 180, 180~
$ drat <dbl> 3.90, 3.90, 3.85, 3.08, 3.15, 2.76, 3.21, 3.69, 3.92, 3.92, 3.92,~
$ wt <dbl> 2.620, 2.875, 2.320, 3.215, 3.440, 3.460, 3.570, 3.190, 3.150, 3.~
$ qsec <dbl> 16.46, 17.02, 18.61, 19.44, 17.02, 20.22, 15.84, 20.00, 22.90, 18~
$ vs <dbl> 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0,~
$ am <dbl> 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0,~
$ gear <dbl> 4, 4, 4, 3, 3, 3, 3, 4, 4, 4, 4, 3, 3, 3, 3, 3, 4, 4, 4, 3, 3,~
$ carb <dbl> 4, 4, 1, 1, 2, 1, 4, 2, 2, 4, 4, 3, 3, 3, 4, 4, 4, 1, 2, 1, 1, 2,~
```

- Compare a select of `mpg` with a pull of `mpg`.

```
mtcars %>% select(mpg) %>% my_gt(6)
```

mpg
21.0
21.0
22.8
21.4
18.7
18.1

```
mtcars %>% pull(mpg)
```



```
[1] 21.0 21.0 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 17.8 16.4 17.3 15.2 10.4
[16] 10.4 14.7 32.4 30.4 33.9 21.5 15.5 15.2 13.3 19.2 27.3 26.0 30.4 15.8 19.7
[31] 15.0 21.4
```

Check the help file and select the second before last column, but `-pull-` it from the data frame so that it turns into a vector

```
mtcars %>% pull(-3)
```

```
[1] 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 0 0 0 0 0 1 1 1 1 1 1 1
```

- Select all columns except the `am`.

```
mtcars %>% select(-am)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	3	3
Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	3	3
Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	3	3
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	3	4
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	3	4
Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	3	4
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	4	1
Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	3	1
Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	3	2
AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	3	2
Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	3	4
Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	3	2

Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	4	1
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	5	2
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	5	2
Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	5	4
Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	5	6
Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	5	8
Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	4	2

- Select all columns except the `am` and `vs`.

```
mtcars %>% select(-am,-vs)
```

	mpg	cyl	disp	hp	drat	wt	qsec	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	3	3
Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	3	3
Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	3	3
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	3	4
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	3	4
Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	3	4
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	4	1
Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	3	1
Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	3	2
AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	3	2
Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	3	4
Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	3	2
Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	4	1
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	5	2
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	5	2

Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	5	4
Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	5	6
Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	5	8
Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	4	2

- Keep only columns mpg, cyl and disp, but rename mpg to miles_gallon.

```
mtcars %>% select(miles_gallon=mpg,cyl,disp)
```

	miles_gallon	cyl	disp
Mazda RX4	21.0	6	160.0
Mazda RX4 Wag	21.0	6	160.0
Datsun 710	22.8	4	108.0
Hornet 4 Drive	21.4	6	258.0
Hornet Sportabout	18.7	8	360.0
Valiant	18.1	6	225.0
Duster 360	14.3	8	360.0
Merc 240D	24.4	4	146.7
Merc 230	22.8	4	140.8
Merc 280	19.2	6	167.6
Merc 280C	17.8	6	167.6
Merc 450SE	16.4	8	275.8
Merc 450SL	17.3	8	275.8
Merc 450SLC	15.2	8	275.8
Cadillac Fleetwood	10.4	8	472.0
Lincoln Continental	10.4	8	460.0
Chrysler Imperial	14.7	8	440.0
Fiat 128	32.4	4	78.7
Honda Civic	30.4	4	75.7
Toyota Corolla	33.9	4	71.1
Toyota Corona	21.5	4	120.1
Dodge Challenger	15.5	8	318.0
AMC Javelin	15.2	8	304.0
Camaro Z28	13.3	8	350.0
Pontiac Firebird	19.2	8	400.0
Fiat X1-9	27.3	4	79.0
Porsche 914-2	26.0	4	120.3
Lotus Europa	30.4	4	95.1
Ford Pantera L	15.8	8	351.0
Ferrari Dino	19.7	6	145.0
Maserati Bora	15.0	8	301.0

```
Volvo 142E                21.4    4 121.0
```

- Using `rename` not all variables need to be mentioned explicitly. Change only `mpg` to `miles_gallon`.

```
mtcars %>% rename(miles_gallon=mpg)
```

	miles_gallon	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

- Keep only the consecutive columns in between `disp` and `wt` (use a `:`), additionally add

mpg as a last column.

- It is possible to pipe also base R functions, try it and pipe the solution above through `names` to get the variable names.

```
mtcars %>% select(dispatch,mpg) %>% names
```

```
[1] "disp" "hp"   "drat" "wt"   "mpg"
```

- Create a variable for the row names. Maybe check `rownames_to_column`.

```
mtcars %>% rownames_to_column('type')
```

	type	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
2	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
3	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
4	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
5	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
6	Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
7	Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
8	Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
9	Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
10	Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
11	Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
12	Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
13	Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
14	Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
15	Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
16	Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
17	Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
18	Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
19	Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
20	Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
21	Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
22	Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
23	AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
24	Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
25	Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
26	Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
27	Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2

28	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
29	Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
30	Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
31	Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
32	Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

- Change the mpg (miles per gallon) into kpl (kilometers per liter) with 1 mpg is 0.425 km/l, using `mutate()`.

```
mtcars %>% mutate(kpl=mpg*.425)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	kpl
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4	8.9250
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4	8.9250
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1	9.6900
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1	9.0950
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2	7.9475
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1	7.6925
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4	6.0775
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2	10.3700
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2	9.6900
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4	8.1600
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4	7.5650
Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3	6.9700
Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3	7.3525
Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3	6.4600
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4	4.4200
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4	4.4200
Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4	6.2475
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1	13.7700
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2	12.9200
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1	14.4075
Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1	9.1375
Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2	6.5875
AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2	6.4600
Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4	5.6525
Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2	8.1600
Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1	11.6025
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2	11.0500
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2	12.9200
Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4	6.7150

Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6	8.3725
Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8	6.3750
Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2	9.0950

- Select about 10% of the observations, check the help file on using `sample_frac()`.
 - run this code multiple times to see what happens

```
mtcars %>% sample_frac(.1)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Dodge Challenger	15.5	8	318	150	2.76	3.520	16.87	0	0	3	2
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1

```
mtcars %>% sample_frac(.1)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Merc 280C	17.8	6	167.6	123	3.92	3.44	18.9	1	0	4	4
Merc 230	22.8	4	140.8	95	3.92	3.15	22.9	1	0	4	2
Porsche 914-2	26.0	4	120.3	91	4.43	2.14	16.7	0	1	5	2

- Select the 10th to 15th row, check the help file on using `slice()`

```
mtcars %>% slice(10:15)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Merc 280	19.2	6	167.6	123	3.92	3.44	18.30	1	0	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.44	18.90	1	0	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.07	17.40	0	0	3	3
Merc 450SL	17.3	8	275.8	180	3.07	3.73	17.60	0	0	3	3
Merc 450SLC	15.2	8	275.8	180	3.07	3.78	18.00	0	0	3	3
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.25	17.98	0	0	3	4

- Select only the distinct combinations, for variables `am` and `vs`

```
mtcars %>% distinct(cyl,vs,am)
```

	cyl	vs	am
Mazda RX4	6	0	1

Datsun 710	4	1	1
Hornet 4 Drive	6	1	0
Hornet Sportabout	8	0	0
Merc 240D	4	1	0
Porsche 914-2	4	0	1
Ford Pantera L	8	0	1

- You only get three variables, check the help files to determine how to keep all variables (for each first observation of that combination)

```
mtcars %>% distinct(cyl,vs,am,.keep_all=T)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4

- Filter the data to retain only cases with mpg > 20 and hp above or equal to 110

```
mtcars %>% filter(mpg>20, hp>=110)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2

- Filter the data to retain only the Datsun 710

```
mtcars %>% rownames_to_column('type') %>% filter(type=='Datsun 710')
```

	type	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	Datsun 710	22.8	4	108	93	3.85	2.32	18.61	1	1	4	1

Getting ahead of ourselves again, with tidier and friends

toy dataset

- A dataset can be read in, for example using the `read_delim()` function
- Just copy-paste data from notepad, excel or another spreadsheet program
- The copy-pasted table can be assigned to the `myrepeated` object

```
myrepeated <- read_delim(clipboard(),delim='\t')
```

- The `clipboard()` function is just one way, you can also specify a path to the data
 - the delimiter is `\t` or TABs
 - type `?read_delim` to get details on more possibilities
- Because you do not have it, it is included already

```
(myrepeated <- tribble(
  ~id, ~`t1 score`, ~`t1 posit`, ~`t2 score`, ~`t2 posit`, ~`t3 score`, ~`t3 posit`,
  "id1",1,'x', NA,'y',4,'x',
  "id2",2,'y',3,'x',NA,NA,
  "id3",1,'x',2,'y',5,'x'
))
```

A tibble: 3 x 7

	id	't1 score'	't1 posit'	't2 score'	't2 posit'	't3 score'	't3 posit'
	<chr>	<dbl>	<chr>	<dbl>	<chr>	<dbl>	<chr>
1	id1	1	x	NA	y	4	x
2	id2	2	y	3	x	NA	<NA>
3	id3	1	x	2	y	5	x

Joining with `'by = join_by(id, time)'`

exemplary data tidying

- Having it read in, it is tidied, turned into 2 files joined after separating cell contents
 - make a dataset without the posit variables, and one without the score variables, and pivot the score or posit values from columns to rows identified by a new variable type

- disentangle the values in **type** in two parts: **time** and **type**
- recombine the two datasets after removing the new variable **type** from at least one of them
- remove all rows with missing values in either the variable **score** or **posit**

```
scores <- myrepeated %>%
  select(id, `t1 score`, `t2 score`, `t3 score`) %>%
  pivot_longer(-id, names_to='type', values_to='score')
positions <- myrepeated %>%
  select(id, `t1 posit`, `t2 posit`, `t3 posit`) %>%
  pivot_longer(-id, names_to='type', values_to='posit')

scores <- scores %>%
  separate(type, c('time', 'type'))
positions <- positions %>%
  separate(type, c('time', 'type'))

joined <- scores %>%
  select(-type) %>% full_join(positions)

longform <- joined %>%
  select(-type) %>% filter(!is.na(score), !is.na(posit))
```

id	time	score	posit
id1	t1	1	x
id1	t3	4	x
id2	t1	2	y
id2	t2	3	x
id3	t1	1	x
id3	t2	2	y
id3	t3	5	x

- It is possible to switch back to a wider data representation
 - e.g., to calculate correlations (maybe fill in the missing values NA as 0 values)

```
longform %>% pivot_wider(names_from=c(time), values_from=c(score, posit), values_fill=list(sc
```

```
# A tibble: 3 x 7
  id      score_t1 score_t3 score_t2 posit_t1 posit_t3 posit_t2
<chr>    <dbl>    <dbl>    <dbl> <chr>    <chr>    <chr>
```

1	id1	1	4	0	x	x	<NA>
2	id2	2	0	3	y	<NA>	x
3	id3	1	5	2	x	x	y

tidyr and import packages, functions to read and tidy data

- **tidier** combines a few functions to tidy up the data
 - a core idea at the origin of the development of the tidyverse
- By enforcing structure on the data, functions defined to operate on that data can be made much more consistent too
- **readr** combines a few functions to read in data, stored externally, in text format, excel, spss, ...
- The **tidier** and **readr** packages:
 - focus on importing data and making it tidy
 - * the data has to be brought into the R workspace
 - * the data has to be tidy for efficient further processing
 - use to create tidy data
 - * a row for each research unit
 - * a columns for each variable
 - * a cell that links a research unit to a variable
 - requires
 - * pivoting data into longer or wider form
 - * creating pure variables
- The main -verbs- (see example above)
 - **pivot_wider()** and **pivot_longer()**: turn multiple columns or rows into one, making datafiles longer or wider
 - **separate()** and **extract()**: create multiple columns from one column using delimiters or regular expressions

pivot_*()

- Turning long form data into wide form and vise versa, is called pivoting.

intro

- In tidy data each research unit is assigned to a row, in a tidy dataframe (tibble)
 - what is the research unit depends on the research question and can change (eg., test score → student)
- Contrary to univariate data representation, a multivariate data representation can be useful and be more intuitive
- To change research units or to switch between uni -and multivariate, data can be pivoted, turned wider or longer
- Pivoting from wider to longer
 - column headers are turned into values of an identifier column
 - values over different columns are combined into new column
 - the identifier column and values column require names
- The `iris` dataset, with 4 values for each unit within each species, is pivoted
 - Notice, the k column headers turn into nxk cell values to serve as identifiers
- Beware: without a unique identifier for each row, the clustering of columns' information is lost
 - a unique identifier per row should typically be added before pivoting
 - given a unique identifier, it should be removed from the pivoted variables

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa

```
long_iris_no_identifier <- iris %>% pivot_longer(-Species, names_to='type', values_to='score')
long_iris_with_identifier <- iris %>% mutate(id=1:n()) %>% pivot_longer(-c(id, Species), names_to='type', values_to='score')
```

Species	type	score
setosa	Sepal.Length	5.1
setosa	Sepal.Width	3.5
setosa	Petal.Length	1.4

setosa	Petal.Width	0.2
setosa	Sepal.Length	4.9
setosa	Sepal.Width	3.0

Species	id	type	score
setosa	1	Sepal.Length	5.1
setosa	1	Sepal.Width	3.5
setosa	1	Petal.Length	1.4
setosa	1	Petal.Width	0.2
setosa	2	Sepal.Length	4.9
setosa	2	Sepal.Width	3.0

- Pivoting from longer to wider
 - column headers are created from values in an identifier column
 - values within a values column are aligned over different columns
 - the identifier column and values column must be specified
- Without adding a row specific identifier before pivoting the `iris` dataset from wide to long
 - no information would be available to assign values to a particular row
 - many values are forced into one single cell
- Note that long-er and wide-r is used
 - expresses that data can be long for certain aspects and wide for others

```
long_iris_no_identifier %>% pivot_wider(values_from=score,names_from=type)
```

Warning: Values from 'score' are not uniquely identified; output will contain list-cols.

* Use 'values_fn = list' to suppress this warning.

* Use 'values_fn = {summary_fun}' to summarise duplicates.

* Use the following dplyr code to identify duplicates.

```
{data} %>%
```

```
dplyr::group_by(Species, type) %>%
```

```
dplyr::summarise(n = dplyr::n(), .groups = "drop") %>%
```

```
dplyr::filter(n > 1L)
```

```
# A tibble: 3 x 5
```

```
Species    Sepal.Length Sepal.Width Petal.Length Petal.Width
<fct>      <list>        <list>      <list>      <list>
```

```

1 setosa      <dbl [50]>    <dbl [50]>    <dbl [50]>    <dbl [50]>
2 versicolor <dbl [50]>    <dbl [50]>    <dbl [50]>    <dbl [50]>
3 virginica   <dbl [50]>    <dbl [50]>    <dbl [50]>    <dbl [50]>

```

- To pivot from longer to wider form, a column is spread out over multiple columns and along with it the values
 - new column names are extracted from a column, typically with a limited set of labels
 - values to populate the newly constructed columns are extracted from a column too

```
long_iris_with_identifier %>% group_by(type) %>% mutate(id=1:n()) %>% pivot_wider(values_f
```

Species	id	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
setosa	1	5.1	3.5	1.4	0.2
setosa	2	4.9	3.0	1.4	0.2
setosa	3	4.7	3.2	1.3	0.2
setosa	4	4.6	3.1	1.5	0.2
setosa	5	5.0	3.6	1.4	0.2
setosa	6	5.4	3.9	1.7	0.4

exercises

- Pivot the `world_bank_pop` dataset from the `tidyr` package, to have univariate data for the scores over the different years

```
world_bank_pop %>% pivot_longer(-c(country,indicator),values_to='scores',names_to='year')
```

```

# A tibble: 19,152 x 4
  country indicator  year  scores
  <chr>    <chr>    <chr> <dbl>
1 ABW     SP.URB.TOTL 2000  41625
2 ABW     SP.URB.TOTL 2001  42025
3 ABW     SP.URB.TOTL 2002  42194
4 ABW     SP.URB.TOTL 2003  42277
5 ABW     SP.URB.TOTL 2004  42317
6 ABW     SP.URB.TOTL 2005  42399
7 ABW     SP.URB.TOTL 2006  42555
8 ABW     SP.URB.TOTL 2007  42729
9 ABW     SP.URB.TOTL 2008  42906
10 ABW    SP.URB.TOTL 2009  43079

```

```
# i 19,142 more rows
```

- Use the `us_rent_income` dataset, also part of the `tidyr` package, and remove variable `moe` before pivoting the estimates to wide form

```
(us_rent_income %>% select(-moe) %>% pivot_wider(values_from=estimate,names_from=variable))
```

```
# A tibble: 52 x 4
```

	GEOID	NAME	income	rent
	<chr>	<chr>	<dbl>	<dbl>
1	01	Alabama	24476	747
2	02	Alaska	32940	1200
3	04	Arizona	27517	972
4	05	Arkansas	23789	709
5	06	California	29454	1358
6	08	Colorado	32401	1125
7	09	Connecticut	35326	1123
8	10	Delaware	31560	1076
9	11	District of Columbia	43198	1424
10	12	Florida	25952	1077

```
# i 42 more rows
```

- Verify what happens when you did not remove `moe`

```
(us_rent_income %>% pivot_wider(values_from=estimate,names_from=variable))
```

```
# A tibble: 104 x 5
```

	GEOID	NAME	moe	income	rent
	<chr>	<chr>	<dbl>	<dbl>	<dbl>
1	01	Alabama	136	24476	NA
2	01	Alabama	3	NA	747
3	02	Alaska	508	32940	NA
4	02	Alaska	13	NA	1200
5	04	Arizona	148	27517	NA
6	04	Arizona	4	NA	972
7	05	Arkansas	165	23789	NA
8	05	Arkansas	5	NA	709
9	06	California	109	29454	NA
10	06	California	3	NA	1358

```
# i 94 more rows
```

- It is possible to include multiple variables to pivot wide, jointly, use a vector of variables that includes estimate and moe, and see what happens

```
(us_rent_income %>% pivot_wider(values_from=c(estimate,moe),names_from=variable))
```

```
# A tibble: 52 x 6
  GEOID NAME          estimate_income estimate_rent moe_income moe_rent
  <chr> <chr>          <dbl>         <dbl>    <dbl>    <dbl>
1 01    Alabama      24476          747      136      3
2 02    Alaska      32940         1200      508     13
3 04    Arizona      27517          972      148      4
4 05    Arkansas      23789          709      165      5
5 06    California    29454         1358      109      3
6 08    Colorado      32401         1125      109      5
7 09    Connecticut    35326         1123      195      5
8 10    Delaware      31560         1076      247     10
9 11    District of Columbia 43198         1424      681     17
10 12    Florida      25952         1077       70      3
# i 42 more rows
```

separate() / unite()

- Splitting up information within a variable, or combining information over variables, to ensure cell values to offer one and only one piece of relevant information

intro

- Each variable should consist of one type of information, in a tidy dataframe (tibble)
 - variables that combine information should often be split
 - variables that provide no meaningful information by themselves should be removed, sometimes united
- Columns (variables) can be split and united
- The long form iris data shows a type that consists of both Petal/Sepal and Length/Width, the can be separated

```
long_iris_with_identifier %>% separate(type,c('PS','lw'))
```


Species	id	PS	lw	score
setosa	1	Sepal	Length	5.1
setosa	1	Sepal	Width	3.5
setosa	1	Petal	Length	1.4
setosa	1	Petal	Width	0.2
setosa	2	Sepal	Length	4.9
setosa	2	Sepal	Width	3.0

- On the contrary, variables can also be united
- Separated columns can be combined, using a separator dash in this case (default is underscore)

```
long_iris_separated %>% unite('myType',PS:lw,sep='-')
```

```
# A tibble: 600 x 4
  Species    id myType    score
  <fct>    <int> <chr>    <dbl>
1 setosa      1 Sepal-Length  5.1
2 setosa      1 Sepal-Width  3.5
3 setosa      1 Petal-Length  1.4
4 setosa      1 Petal-Width  0.2
5 setosa      2 Sepal-Length  4.9
6 setosa      2 Sepal-Width   3
7 setosa      2 Petal-Length  1.4
8 setosa      2 Petal-Width  0.2
9 setosa      3 Sepal-Length  4.7
10 setosa     3 Sepal-Width  3.2
# i 590 more rows
```

- The `tidyr` package includes other functions for more involved programming and simulation studies
- Notice in particular `expand()`, `crossover()`, `nesting()`, best check the helpfile.

[?expand](#)

exercises

- Turn the row names of the `mtcars` data to a variable called `type` using the `rownames_to_column()` function

- it consists of car type information (maybe use `type`), car sub-type (`subtype`) and sub-type specification (`spec`)
- look into the `fill` argument to ensure the the pieces of information are read in from right to left

```
mtcars %>% rownames_to_column('type') %>%
  separate(type,c("type","subtype","spec"),fill='right') %>%
  ungroup() %>% my_gt(6)
```

type	subtype	spec	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda	RX4	NA	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda	RX4	Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun	710	NA	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet	4	Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet	Sportabout	NA	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	NA	NA	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

- Separate the type variable to isolate information on the type on one hand, and the rest on the other

```
mtcars %>% rownames_to_column('type') %>%
  separate(type,c("type","subtype","spec"),fill='right') %>%
  unite("subtype",c("subtype","spec")) %>% ungroup() %>% my_gt(6)
```

type	subtype	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda	RX4_NA	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda	RX4_Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun	710_NA	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet	4_Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet	Sportabout_NA	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	NA_NA	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

Import data with `readr`, `readxl` or `haven`

- when using your own data, they have to be imported into the workspace
- Data that are saved as R objects in a workspace (`*.RData`) can be loaded with the `load()` function
- Data that need to be imported from elsewhere require dedicated functions (packages)

readr

- The `readr` package in tidyverse deals with the basic data, like comma separated or tab-delimited data

intro

- The primary function in `readr` is `read_delim()` which imports tabular data with a delimiter as specified
- Note that a path to the data may need to be specified, in absolute terms or relative to the current working directory

```
getwd()
setwd(readClipboard())
setwd('../.. /my_sub_dir_2_levels_up')
```

- A delimiter should be specified, `\t` for tabs
- `?read_delim` offers information on how to set many different arguments and gain flexibility to read in data
- In this current working directory should have a tab-delimited file named `repeated.txt`

```
myrepeated <- read_delim(file='repeated.txt',delim='\t') # if
```

- Data can be copy pasted in using the `clipboard()` instead of a path, or a path can be asked for interactively with `file.choose()`

```
myrepeated <- read_delim(clipboard(),delim='\t')
myrepeated <- read_delim(file.choose(),delim='\t')
```

readxl

- The `readxl` package in tidyverse deals with the notorious excel files

intro

- The primary function in `readxl` is `read_excel()` which imports tabular data from an excel file
- Note that a path to the data may need to be specified, in absolute terms or relative to the current working directory
- The `example_data_set.xlsx` if it would exist in current working directory could be read in, possibly having assigned a particular sheet
- Interesting arguments are the sheet to read from, or the number of rows to skip
- `?read_excel` offers information on the many arguments that add flexibility for reading in data

```
read_excel('example_data_set.xlsx', sheet='my_data', skip=1)
```

haven

- The `haven` package in tidyverse deals with the data stored as part of one of the main statistical software, like SAS, spss and Stata

intro

- For SPSS, with *.sav files, Data is simply read, using default parameters `read_sav()` reads SPSS stored data
- The `haven` package is not automatically loaded with tidyverse
- Let's first get the path to the iris data as an example
- `?read_sav` for more information on the available arguments

```
library(haven)
path_to_spss_exemplary_data <- system.file("examples", "iris.sav", package = "haven")
read_sav(path_to_spss_exemplary_data)
```

A tibble: 150 x 5

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl+lbl>
1	5.1	3.5	1.4	0.2	1 [setosa]
2	4.9	3	1.4	0.2	1 [setosa]
3	4.7	3.2	1.3	0.2	1 [setosa]

```

4          4.6          3.1          1.5          0.2 1 [setosa]
5          5           3.6          1.4          0.2 1 [setosa]
6          5.4          3.9          1.7          0.4 1 [setosa]
7          4.6          3.4          1.4          0.3 1 [setosa]
8          5           3.4          1.5          0.2 1 [setosa]
9          4.4          2.9          1.4          0.2 1 [setosa]
10         4.9          3.1          1.5          0.1 1 [setosa]
# i 140 more rows

```

- For SAS, with for example *.sas7bdat files, data is read using default parameters
- Let's again get the path to the iris data as an example
- `?read_sas` for more information on the available arguments

```

path <- system.file("examples", "iris.sas7bdat", package = "haven")
read_sas(path)

```

- For Stata, with for example *.dta files, data is read using default parameters
- Let's again get the pat to the iris data as an example
- `?read_dta` for more information on the available arguments

```

path <- system.file("examples", "iris.dta", package = "haven")
read_dta(path)

```

To write any of the files, use the `write_` prefix, for `dta`, `sas` and `sav`
 To write the `mtcars` into sas format.

```


write_sas(mtcars, 'mytryinSAS.sas7bdat')

```

Last remarks

Current page provides a primer on data manipulation, tidying data and the importing of data, which are the main steps in preparation of most real data analyses and visualizations.

It is strongly advised to play with the techniques discussed above to get some proficiency in using it, as it would add significantly to the flexibility of whatever you want to further do with your data.



Other tidyverse packages exist, and within the same framework many more are being developed. The consistency within the tidyverse ecosystem should give you a push though, to study the other packages yourself when of interest.

Base R still is a proper alternative to the tidyverse ecosystem, so be aware that others may do things differently.